

# Evaluativity

**Isidora Stojanovic**

Institut Jean-Nicod (CNRS & ENS, PSL Research University),  
29 rue d'Ulm, 75005 Paris, France

Chapter draft (August 2022) for the Oxford Handbook of Contemporary  
Philosophy of Language, ed. Ernie Lepore & Una Stojnic

**Keywords:** evaluative adjectives; value; valence; polarity; judge-dependence; subjectivity

## 1 Introduction

The notion of *evaluativity* may be understood in many different ways. One dimension of variation concerns its bearers. Words, terms, expressions, constructions, predicates, properties, sentences, propositions, statements, or judgments, can all be characterized as "evaluative". The different notions of evaluativity that vary across this dimension are often interestingly intertwined. For example, a linguistic construction that contains an evaluative word will, under appropriate conditions, be itself evaluative, and, if uttered in appropriate conditions, will result in an evaluative statement, which will express an evaluative judgment. Now, there is another, more interesting divide in theorizing about evaluativity. There is one theoretical approach that links evaluativity tightly to the notion of (positive or negative) *valence*. In a nutshell, according to this approach, to say about someone that they are beautiful, good, or generous, is to say something positive about them, whereas to say that they are ugly, bad, or selfish, is to say something negative. Relatedly, adjectives 'beautiful/ugly', 'good/bad', 'generous/selfish', are considered to be evaluative adjectives *par excellence*. And then, there is another theoretical approach that links evaluativity tightly to the notion of *value*. While we can, of course, speak of *positive* and *negative* value, thereby incorporating valence into a value-system, the notion of value provides a different framework to think about evaluativity. Evaluative expressions are, then, those whose function is to attribute values. 'Beautiful' and 'ugly' come out, again, as evaluative, because they serve to rank objects along a scale of aesthetic value. But even expressions such as 'expensive' and 'cheap' can be seen as "evaluative" according to this view, since they rank objects according to their commercial value. Scales like commercial value, however, need not be polarized or reflect any kind of valence.

Our main goal in this chapter is to articulate these two notions of evaluativity. While both notions have been central in the literature on evaluative language, the distinction between them has not always been carefully drawn, and their relationship

has hardly been discussed at all. We hope to remedy this shortcoming. Also, let us note from the outset that there is yet a third notion of evaluativity, fairly prominent in the linguistic literature, according to which “an adjectival construction is evaluative if and only if it conveys that the property associated with the adjective exceeds a relevant threshold” (Brasoveanu & Rett 2017, p. 1; see also Bierwisch 1989 and Rett 2014). Even if this notion has some connection with evaluativity understood in terms of attribution of value, we shall mostly leave it aside.

This chapter is organized as follows. Section 2 explores the Valence Approach to evaluativity; Section 3, the Value Approach. Section 4 presents several linguistic criteria for identifying evaluative terms, and compares evaluativity to the closely related notion of subjectivity. Section 5 outlines further research issues and concludes.

## 2 Understanding evaluativity in terms of VALENCE

In philosophy, the notion of evaluativity has been traditionally tied to the notion of positive or negative *valence* (or positive/negative *polarity*).<sup>1</sup> In his influential book on evaluative language, Pekka Väyrynen characterizes evaluative terms as those that convey “information to the effect that something has a positive or negative standing—merit or demerit, worth or unworth—relative to a certain kind of standard” (Väyrynen 2013, p. 29). Historically, the focus on valence has been a core tenet of expressivism about evaluative (and, in particular, moral) language. As Elisabeth Camp puts it, “according to expressivists, when I say “Murder is wrong!” I don’t describe a state of affairs, but avow or display or advocate a *negative* attitude toward murder” (Camp 2017, p. 87; our italics). As Mark van Roojen puts it, “expressivists of all sorts think that moral sentences are conventional devices for expressing pro and con attitudes towards their objects” (van Roojen 2018, sect. 4.1). The very idea of *pro* and *con* attitudes mirrors the binary division between positive and negative valence. Positive evaluative terms, statements and judgments invite pro attitudes, such as approval and support, while negative ones invite con attitudes, such as disapproval and rejection.

The focus on valence leads to a simple and elegant picture of evaluativity, which can be summarized as follows:

- Evaluative words, terms, or expressions, are those whose meaning encodes either a positive valence ('good', 'excellent', 'beautiful', 'generous', 'attractive') or a negative valence ('bad', 'horrible', 'ugly', 'selfish', 'repulsive').
- Evaluative sentences are those that, when asserted sincerely, give rise to evaluative statements, which in turn, express evaluative judgments. A judgment is

---

<sup>1</sup> While both terms are widely used in the literature, we prefer to use 'valence' for at least two reasons. First, we wish to avoid any confusion with so called *negative polarity items*, such as 'any' or 'ever'. Second, polarity is arguably a broader notion: “Polarity is, in essence, the relation between semantic opposites - between meanings (or expressions denoting meanings) which are fundamentally inconsistent with each other” (Israel 2004, p. 701). Thus the antonyms *hot-cold*, *long-short*, *simple-complex* are polar opposites, but they are not valenced.

evaluative if it depicts something in either a positive or a negative light.<sup>2</sup>

- Typically, an evaluative term, used in the predicate position of an affirmative sentence, will give rise to an evaluative statement; and, conversely, an evaluative statement is typically formulated with the help of evaluative terms.

While clearly attractive, this picture is not without problems, of which we tackle only a few in the remainder of this section.

## 2.1. How can we tell whether a word has a (positive or negative) valence?

In the examples given so far -- good vs. bad, excellent vs. horrible -- the corresponding valence appears to be constitutive of the word's meaning. Consider a person who sincerely uses the word 'good' to describe things about which they feel negatively, and 'bad' to describe those about which they feel positively. We would consider such a person to be an incompetent speaker of English. However, once we move away from such basic, all-purpose evaluatives, determining a word's valence is no longer a trivial matter. To see why, let us take a look at how valence is understood in linguistics.

When linguists speak of valence, they typically speak of *affective* meaning; that is, subjective emotions and attitudes that speakers and hearers associate with words (or linguistic constructions more generally). Word, and generally constructions, inherit the valence of the associated emotional response. Psycholinguists have conducted studies over large numbers of participants to elicit judgments concerning the valence of thousands of words (see i.a. Warriner 2013, Mohammad 2018). When we look at the results, we see that participants' judgments often largely converge, in which case it makes sense to consider the word itself as having a lexically encoded, stable valence. Nevertheless, we must bear two points in mind. First, psychological norms of valence target specifically the participants' emotional responses. Second, participants are not asked about a binary distinction between positive and negative words, but rather, they are asked to rate a word on a continuous *scale*. Here is an example of the instructions that are given to participants:

"You are invited to take part in the study that is investigating emotion, and concerns how people respond to different types of words. You will use a scale to rate how you felt while reading each word. The scale ranges from 1 (happy) to 9 (unhappy). At one extreme of this scale, you are happy, pleased, satisfied, contented, hopeful. (...) The other end of the scale is when you feel completely unhappy, annoyed, unsatisfied, melancholic, despaired, or bored. (...) If you feel completely neutral, neither happy

---

<sup>2</sup> This admittedly vague formulation aims to stay neutral with respect to the cognitivism-non-cognitivism divide. That is to say, it is meant to be compatible with the view that an evaluative judgment expresses a doxastic attitude (belief, knowledge) the content of which incorporates the corresponding positive or negative valence, as well as with the view that evaluative judgments express non-doxastic attitudes such as (dis)approval, endorsement or rejection.

nor sad, select the middle of the scale" (Warriner et al. 2013, p. 1193).<sup>3</sup>

This raises several problems. First, terms relating to illness or death end up being among the most negatively valenced ones. Thus in Warriner et al 2013, in the top five of negative words we find AIDS (1.33) and leukemia (1.47). To be sure, saying of someone that they are diagnosed with AIDS, or describing some substance as carcinogenic, evokes deeply negative feelings; however, such statements lie far away from what philosophers consider to be evaluative statements. To the contrary, stating that someone has AIDS, or that a substance is carcinogenic, could even be seen as paradigmatically descriptive, factual statements.

More generally, while the psychologists' notion of valence is tightly linked to feelings and emotions, the philosophers' notion of evaluativity aims to track values rather than emotions (see section 3). In the metaethical literature on evaluative terms, it is often taken for granted that even terms such as 'pleasant' and 'painful' are *not* evaluative.<sup>4</sup> Yet, from a more psychological perspective, when we describe something as painful, we normally convey a strongly negative evaluation about the thing at stake.

The second problem is that words are presented to participants out of context, which means that they are likely to rate the word for valence according to the feeling that would be typically evoked by the use of this word. This obliterates the fact that the valence of some words may vary with the context. Consider the adjective 'intense'. Whether by describing something as "intense" one will convey something positive or negative will typically depend on what one is talking about. If one is describing the pain they felt after an injury, then its intensity will likely be a bad thing, whereas if they are describing the pleasure that they felt on some occasion, it becomes a good thing.<sup>5</sup> Even terms with which we normally associate negative emotions, such as 'terrifying', 'scary', 'harrowing' or 'disturbing', sometimes convey positive evaluations.

---

<sup>3</sup> Psychological norms are also used to elicit participants' judgments regarding two other dimensions: *arousal* and *dominance*. For simplicity, we will ignore those in the present chapter, which is why we have deleted reference to those in the cited passage.

<sup>4</sup> See e.g. Väyrynen 2013, pp. 10, 37, 137, 181-2, 196, 239-40. Väyrynen does not explain *why* terms such as 'pleasant' and 'painful' are not considered to be evaluative, and his own theory is compatible with taking these terms to be evaluative to the same extent as terms like 'generous', 'selfish', 'courageous' or 'coward'. Note that, in contrast with terms like 'carcinogenic', terms such as 'painful' and 'pleasant' are seldom taken to be descriptive either. The reason is that they are deeply subjective, and, as such, closely connected with evaluative terms (see section 4.4.)

<sup>5</sup> It is not, then, surprising that in Warriner et al. (2013), 'intense' has a mean rating of 5.65. For a discussion about how 'intense' can acquire positive or negative valence depending on the context, see Stojanovic (2016). See also Stojanovic and Kaiser (2022) for a discussion of a broader range of *intense*-like adjectives, such as 'excessive' or 'extraordinary'.

This happens, for instance, when we use them in reference to certain works of art.<sup>6</sup>

Thirdly, valence is *prima facie* a binary notion: we talk of positive vs. negative valence. However, psychological norms yield a graded notion, given that participants are asked to rate words for valence on a 1 to 9 point scale. Let us look more closely at this discrepancy in the next section.

## 2.2. Does valence come in degree?

We have just observed that there is a certain tension between the notion of valence that emerges from psychological norms, which is a graded notion, and the pretheoretic notion of valence, which is binary. According to the former, words are not positive or negative *tout court*, since some are more positive/negative than others. According to the latter, words split into three classes: positive, negative, and those that do not carry any valence. The binary notion of valence resonates well with expressivism about evaluative language: positive evaluative statements invite *pro* attitudes, such as endorsement and recommendation, while negative ones invite *con* attitudes, such as rejection and avoidance.

Can the graded notion of valence that comes from psychological norms be put in correspondence with the binary one? This is not immediately clear. Consider 'difficult' and 'easy', whose respective valence ratings reported in Warriner et al. (2018) are 3.78 and 7.47. Considered as a pair of antonyms, 'difficult' is negative while 'easy' is positive. But does this mean that 'difficult' is outright negative while 'easy' is outright positive? This is less clear. Stojanovic and Kaiser (2022) propose several criteria to distinguish *neutral* adjectives from *evaluative* adjectives (see Section 4), and argue that 'difficult' and 'easy' are more appropriately classified as neutral than evaluative. The reason is, roughly, that there are many instances of 'difficult' that are positive (in particular, in relation to games and puzzles) and conversely, many instances of 'easy' that are not positive. Similar observations apply to the pair of antonyms 'simple' and 'complex'. More generally, if we are simply looking at a continuous scale from 1 to 9, it will be very implausible to try to posit any cut-off point to delineate positive words.

Now, it could also be that what we have presented as a potential problem for the graded notion of valence is more of an advantage than a disadvantage. For, it could be that the binary notion of valence is, ultimately, an extrapolation over the graded one. The continuous valence scale does, after all, have a direction: its lower end is negative, its higher end, positive. This is compatible with the idea that valence *is* graded; or, in other words, that it is a property that comes in degree. After all, we do speak of certain things being *more* positive, or *more* negative, than others, and we do

ex, expect to find  
naturalistic  
differences in use  
from corpus data

this can maybe be  
done statistically tho,  
if searching for  
maxima..

---

<sup>6</sup> In philosophy of art, there is a related but separate discussion on how negative emotions contribute positively to aesthetic appreciations; for references, see the section "The Paradox of Negative Emotions" in Contesi (2018). From a more language-oriented point of view, Stojanovic (2016) discusses how the valence of these predicates can shift from negative to positive in contexts of aesthetic evaluation. Cepollaro (2017) offers an account of valence reversals in evaluative predicates that builds on the relevance-theoretic account of slur appropriation from Bianchi (2014).

compare things in terms of *how* good they are, some being *better* than others.

In the next section, we shall explore further the idea that evaluation, rather than being a yes-or-no matter (positive vs. negative), is, at the bottom, a matter of ranking objects, individuals, events, states of affairs, etc., according to their value. For the time being, let us point out that if one settles on the binary notion of valence, it is not obvious how one could account for those phenomena that appear to require a graded notion, such as, in particular, comparative constructions involving evaluatives.<sup>7</sup>

### 3 Understanding evaluativity in terms of VALUE

In the previous section, we saw how the notion of evaluativity (regardless of whether we apply it to terms, statements or judgments) can be plausibly cashed out in terms of the notion of positive and negative *valence*, and relatedly, pro and con attitudes. However, if one focuses exclusively on valence, one loses sight of another crucial notion that lies behind our intuitive understanding of evaluativity, namely, the notion of *value*. Admittedly, this notion, too, leaves room for many different interpretations. One's understanding of value will vary hugely depending on the perspective from which one looks at it: ethics, economy, aesthetics, practical deliberation, what not. Fortunately, we do not have to make any substantive assumptions, or even venture into value theory, in order to see how this notion can lead to another theoretical approach to evaluativity.

Leaving, then, the notion of value as underspecified as possible, there are two widely accepted philosophical ideas that come with this approach. The first is the fact-value distinction; the second, the idea that values come in degrees. Let us look briefly at each.

#### 3.1. 'Evaluative' in opposition to 'descriptive'

When we speak and think about evaluative language, it is customary to contrast it with descriptive language: the former expresses values, the latter states facts. The contrast may be illustrated with these three statements about the Korean film *The*

---

<sup>7</sup> Soria Ruiz (2019) articulates the challenge from evaluative comparatives for expressivist approaches to evaluative discourse in what he dubs "the subsentential Frege-Geach problem". In a nutshell, just as the standard Frege-Geach problem challenges expressivism to explain how constructions involving evaluatives can figure in truth-functional composition ("If A is good, then B"), the challenge raised by the subsentential variant is to explain how 'good' compositionally contributes to 'better (than)' and 'best'. Soria Ruiz (2019) also outlines a semantic account on behalf of expressivism that aims to address this challenge.

*Wailing (Goksung)*:<sup>8</sup>

1. This is another superb film from Na Hong-jin.
2. Terrible. Overly-long, awfully executed film.
3. The film opens with a police sergeant, Jong-gu, investigating a double murder in his quiet, rural village.

The first two statements express aesthetic evaluations of the film at stake, while the third simply describes the opening of the film. The evaluations in these examples are clearly valenced: the first is positive, the second, negative. This should not come as a surprise: we expect the paradigmatic examples of evaluative language to come out as "evaluative" as much on the Valence Approach as on the Value Approach. But the idea that there are evaluative, as opposed to descriptive, statements, is meant to be broader than applying just to clearly positive or negative evaluations. Consider this statement, once again about the film *The Wailing*:<sup>9</sup>

4. Na Hong-jin (...) creates a visceral, moving film that is simultaneously hideous and exquisite, heart-breaking and horrifying.

Intuitively, (4) should count as an evaluative statement, and, relatedly, adjectives such as "visceral", "moving", "hideous", "exquisite", "heart-breaking" and "horrifying" should belong among evaluative adjectives. The statement in (4) reflects the speaker's stance toward the film, their opinion and perception of the film, as opposed to conveying factual information about the film, such as the fact that it is a Korean film, or that it lasts two and half hours. Nevertheless, whether this stance or opinion is necessarily positive or negative is not even clear, and one's inclination to consider (4) as a positive evaluation of the film is largely driven by pragmatic considerations about the writer's choice of words and the features that are usually appreciated for this genre of films.

If valence is no longer our main guide to evaluative language, how do we distinguish it from descriptive, factual language? Two lines of thought may be found in the literature. The first is the idea that evaluative matters cannot be settled by looking at what the world is like, that is, by looking merely at facts. This idea is mirrored by the possibility of persistent disagreements on matters of value. Is Frida Kahlo a greater artist than Henri Matisse? Under which circumstances (if any) are euthanasia and abortion morally permissible? Is turning in an overdue review more important than attending a friend's birthday party? The observation that disagreements in such cases tend to persist even when the disagreeing parties have been informed of

---

<sup>8</sup> Whenever appropriate, we try to provide corpus examples. (1) comes from Jonathan Hathfull, SciFiNow, 20 July 2016: <https://www.scifinow.co.uk/reviews/the-wailing-film-review-fantasia-2016/>; (2) comes from Sam Kearny, Rotten Tomatoes Audience Reviews, 19 November 2016: [https://www.rottentomatoes.com/m/the\\_wailing/reviews/?type=user](https://www.rottentomatoes.com/m/the_wailing/reviews/?type=user); (3) comes from Glenn Kenny, New York Times, 2 June 2016: <http://www.nytimes.com/2016/06/03/films/the-wailing-na-hong-jin-review.html>

<sup>9</sup> Brent McKnight, <https://www.thelastthingisee.com/2016/06/the-wailing-2016-film-review.html>

all the relevant facts has given rise to a lively debate at the interface of philosophy of language and epistemology. Closely related is the idea that there can be so called *faultless* disagreements over matters of personal taste. It is important, however, to keep disagreements over values separate from disagreements over (clearly) subjective matters (see Stojanovic 2019a; see also section 4.4.). For the time being, let us maintain the idea that evaluative and factual statements differ with respect to the speaker's justification for making the statement. While factual statements rely on criteria such as knowledge, evaluative statements appeal to systems of values and norms endorsed by the speaker and their interlocutors.

The second line of thought that may help distinguish values from facts, and relatedly, evaluative from descriptive language, is the notion of action-guidance. The idea, in a nutshell, is that evaluative judgments motivate action in a way in which factual judgments do not. This idea has been particularly influential in metaethics, and has played a crucial role in the debate between cognitivism and non-cognitivism (see Bedke 2017; van Roojen 2018; Soria Ruiz 2019; see also Soria Ruiz, Cepollaro & Stojanovic 2021 for action-guidance regarding value judgments more generally).

For our purposes, let us briefly reflect a difference between statements such as (1) or (2), and (3). If you are told that *The Wailing* is a superb film, and you trust this judgment, you will be motivated to go and see it yourself. If you are told that it is a terrible film, and you trust this judgment, then you will do the opposite. But if you are told how the film begins, or you are simply told that it is a Korean film, this will not lead you on its own to act in any particular way. For a statement such as (3) to motivate action, there must be other action-guiding judgments that serve as a bridge, such as, say, the judgment that a thriller that begins in the way described promises to be good (or, conversely, boring).

### 3.2. Evaluation as assigning values (along a scale)

Intuitively, to evaluate something is to assign it a value on some relevant scale. To illustrate the idea, suppose that you are asked to evaluate a research project. You are typically given a number of criteria, such as originality, soundness, timeliness, and so on. Each criterion typically comes with a scale, such as <poor, average, good, very good, excellent, outstanding>, or simply a numerical scale. You assign the project a value on each such scale. Some criteria may have more weight than others. In the end, based on the different criteria's weights and the values that you have picked out for each, you compute the overall value.

Often, pairing an object with a value on a given scale is a step toward evaluation in a broader sense, since value assignment often only leads to establishing a *ranking* among objects. To continue with our example, consider a project whose score is 7 out of 10. To decide how good it really is, one will want to know more; specifically, one will want to know how its score compares to the scores assigned to other projects. If most other projects get a score in the 8-9 range, then this project is not *that good*; if their average score is, rather, in the 4-5 range, then it really is good.

Now consider the following statements:



5. Béatrice Dalle is older than Sophie Marceau.
6. Béatrice Dalle is more beautiful than Sophie Marceau.
7. Béatrice Dalle is old.
8. Béatrice Dalle is beautiful.

If we pursue the idea that evaluation consists in establishing a ranking along various scales, then (5) throughout (8) would qualify for the status of evaluative statements. At the same time, (5) and (7) seem to be interestingly different from (6) and (8); the former are about age, and thus make reference to factual matters, while the latter, being about beauty, make reference to (aesthetic) value. Note that a statement such as (5) does not give rise to any dispute in the way in which (6) might. One who knows the respective dates of birth of the two actresses cannot reasonably dispute that (5) expresses a truth, whereas (6) can be reasonably disputed by anyone who finds Sophie Marceau at least as beautiful as Béatrice Dalle.

Note also that (7), unlike (5), can generate a disagreement that cannot be settled merely by appealing to facts. This is because in order to decide on a truth value for (7), one must first decide from which age on to consider people "old". Such decisions often depend on the context. A person aged 57 (Dalle's actual age) may be considered old relative to a context in which we are talking about millennials, but not considered old with respect to a discussion about social benefits for senior citizens.

To bring the point home, if we think of evaluation as value assignment along a scale, we get a very broad definition of evaluativity. The latter can be refined, though, in various ways. One can hold that all and only those statements whose truth cannot be settled simply by looking at facts are to be classified as evaluative, thereby excluding (5). Or one can hold that all and only those constructions for which the relevant scale is polarized (as is, presumably, the scale of beauty) ought to qualify as evaluative, thereby excluding both (5) and (7).

One could also hold that only those constructions that convey either a positive or a negative evaluation are to be classified as evaluative, thereby arguably excluding not only (5) and (7), but even (6). The thought would be that a statement such as (6) simply *compares* the two actresses in terms of beauty, without implying that either of them *is* beautiful, hence, does not necessarily say anything positive (or negative) about either of them.<sup>10</sup> To be sure, in a suitable context, (6) *can* convey a positive evaluation about Béatrice Dalle; for instance, if it is taken for granted that Sophie Marceau *is* beautiful. But so can (5) and (7): it can happen that the speaker and their interlocutors share a negative (or, conversely, a positive) attitude toward old people, in which case (7) will convey a corresponding evaluation (and so will (5), on the

---

<sup>10</sup> Adjectives such that their comparative constructions do not entail anything about either of the two relata of comparison are called *relative* gradable adjectives (see i.a. Kennedy & McNally 2005). 'Young' and 'old' are such: from the statement that Sophie Marceau is younger than Béatrice Dalle one cannot infer whether either of them is young or old. Adjectives whose comparative constructions do entail something about one of the two relata are called *absolute*. Whether 'beautiful' and other aesthetic adjectives are relative or absolute is a controversial issue; see Liao and Meskin (2017), Liao, McNally and Meskin (2016), and Stojanovic (2019b) for discussion.

assumption that Sophie Marceau is old.) This suggests, in turn, that one can refine the notion of evaluativity in two ways. The first is to consider only those linguistic constructions that necessarily convey either a positive or a negative evaluation as "evaluative". The second is not to take linguistic constructions to be evaluative *per se*, but rather, to only consider their uses in particular contexts as "evaluative".

Finally, the linguists' notion of evaluativity that we mentioned in Section 1, according to which adjectival constructions that convey that the associated property exceeds a relevant threshold are to be classified as evaluative, could also be seen as a possible refinement of the Value Approach to evaluativity, predicting that (7) and (8), unlike (5) and (6),<sup>11</sup> are evaluative.

## 4 Linguistic criteria for evaluativity

Depending on whether one tackles the notion of evaluativity along the lines of the Valence Approach or the Value Approach, one will end up identifying different, albeit overlapping classes of evaluative expressions. Relatedly, there will be differences as to which statements and judgments are predicted to be evaluative. At the same time, the paradigmatic cases of evaluativity, such as discourse about what is good and bad, what is beautiful and ugly, what is generous and selfish, will end up in the evaluative class on either approach. While we have tried to explain, in the previous two sections, the driving ideas behind the two approaches, there remains the issue of how to identify evaluativity in practice. Let us presently narrow the focus of our inquiry to evaluativity as it applies to words, terms and expressions. One can then try to come up with linguistic criteria that determine whether a given expression is evaluative or not.

The methodology of linguistic criteria, or *tests*, is a standard methodological practice in semantics, and, in particular, in adjectival typology. Just to mention a few examples, McNally and Kennedy (2005) devise tests to classify adjectives between gradable and non-gradable and, among the former, between relative and absolute. Sæbø (2009) offers a test involving embeddings under the verb 'find' to identify subjective adjectives (see section 4.4.). Sassoon (2013) proposes tests to identify multidimensional adjectives. McNally and Stojanovic (2017) use tests to tease apart predicates of personal taste from aesthetic adjectives, and Stojanovic (2019a), to tease them apart from moral adjectives. Our aim in this section is to see how the test methodology can be pursued in theorizing about evaluativity. We will look at six tentative tests to identify evaluative terms. In section 4.1., we will look at the two tests proposed in Stojanovic and Kaiser (2022), which rely on the Valence Approach. In discussing them, we will also forestall certain worries with the test methodology in general. In sections 4.2. and 4.3., we will look at four tests proposed in Soria Ruiz and Stojanovic (2019), which rely on the Value Approach: two that exploit the idea that

---

<sup>11</sup> Once again, though, if 'beautiful' is taken to be an absolute rather than relative gradable adjective, then comparative constructions containing 'beautiful' will also be evaluative, because they will convey that the first relatum exceeds the relevant threshold. See also Bierwisch (1989) on the antonyms 'pretty' and 'ugly'.

evaluative terms are action-guiding, and two that appeal to differences in epistemic justification. In section 4.4., we will turn to two tests that have been widely discussed in the literature, namely, faultless disagreements and embeddings under subjective attitudes such as 'find'. However, we will suggest that those two tests track *subjectivity* rather than *evaluativity*; see Willer (this volume).

#### 4.1. Two valence-based tests for evaluativity

In their recent article "Exploring Valence in Judgments of Taste", Stojanovic and Kaiser are interested in what they call *neutral* predicates of personal taste: adjectives such as 'simple', 'surprising' and 'intense', which are subjective but, unlike the usual predicates of personal taste, such as 'delicious', 'fun', 'disgusting' and 'boring', are neither clearly positive nor clearly negative. Nevertheless, such neutral predicates can, and often do, convey evaluations, in the sense of conveying something positive or negative about the object to which they are ascribed. However, neutral predicates acquire valence only in a context, while evaluative predicates carry valence as part of their lexical meaning.

Stojanovic and Kaiser advocate a two-pronged approach to the topic of neutral predicates. On the one hand, they use psychological norms, of the sort discussed in section 2.1, to identify adjectives whose valence is neither clearly positive nor negative but oscillates around the middle of the associated valence scale. On the other hand, they propose two linguistic tests to distinguish neutral and evaluative adjectives. In combination, the two criteria - norms and tests - reveal a rich class of hitherto neglected adjectives, which are not lexically valenced but can convey evaluations in context. These include: *surprising, strange, intriguing, lavish, intense, difficult, easy, complex* and *simple*.

The first linguistic test to distinguish evaluative and neutral adjectives is the Attitude Compatibility Test. The idea behind it is that if an adjective is clearly positive, and if this information is encoded in its meaning, then this leads to the expectation that one's attitude toward the object to which the adjective is ascribed should also be positive; *mutatis mutandis* for the negative ones. In turn, ascribing a positive adjective to something while at the same time denying that one has such a positive attitude should sound marked, and similarly for ascribing a negative adjective while asserting one's positive attitude. No such markedness, on the other hand, should arise for neutral adjectives. To see how the test works, compare the following minimal

pairs:<sup>12</sup>

- 9. ?? This game is fun and I don't like it.
- 10. ?? This game is boring and I like it.
- 11. ✓ This game is surprising and I (don't) like it.
- 12. ✓ This game is intense and I (don't) like it.

To be sure, from a logical point of view, there is no incompatibility in not liking something that has something positive to be said about it, nor is there in liking something that has some negative feature. Crucially, however, when the valence of the ascribed adjective does not align with the valence of the reported attitude, the speaker is expected to use contrastive conjunctions such as *but* instead of *and*:

- 13. ✓ This game is fun but I don't like it.
- 14. ✓ This game is boring but I like it.

The second linguistic test introduced by Stojanovic and Kaiser is the Good/Bad Way Test. It relies on the idea that to have a property that is neither clearly positive nor negative can sometimes be a good thing, and sometimes a bad thing, depending on the context. No such variation, on the other hand, occurs (unless exceptionally) with evaluative properties. Again, observe the following contrasts:

- 15. ?? This game is fun in a good way.
- 16. ?? This game is fun in a bad way.
- 17. ?? This game is boring in a good way.
- 18. ?? This game is boring in a bad way.
- 19. ✓ This game is surprising/intense/difficult in a good way.
- 20. ✓ This game is surprising/intense/difficult in a bad way.

In the case of (15) and (18), the oddity arises from the fact that the adjective already encodes the corresponding valence. The statements sound odd because it is redundant to say of something positive that it is so "in a good way", and conversely, to say of something negative that it is so "in a bad way". As for (16) and (17), they sound odd because it is hard to think how something positive can be so in a bad way, and how

---

<sup>12</sup> The examples are from Stojanovic and Kaiser (2022), who discuss neutral adjectives among predicates of personal taste. How a contrast between evaluative and neutral adjectives can be drawn in other domains is not completely straightforward. Consider moral adjectives. There is a certain markedness in saying "She is generous and I don't like her" or "She is selfish and I like her". (To be sure, one can dislike generous persons and like selfish ones, but to state this, we would expect the speaker to use a contrastive *but* instead of *and*.) However, to show that there is no such markedness for putatively neutral moral adjectives presupposes that we can delineate the class of moral adjectives in the first place. For predicates of personal taste, this can be easily done (Stojanovic and Kaiser rely on a set of criteria for identifying PPTs, largely following Umbach (2021); see also Willer (this volume) and section 4.4. below). When it comes to moral predicates, one could think of personality traits, such as *extravagant*, *meticulous* or *ambitious*, all three of which are neither clearly positive nor clearly negative, and which clearly pass the Attitude Compatibility Test. However, it can be disputed whether such adjectives ought to be classified as "moral" adjectives.

something negative can be so in a good way.

It bears noting, however, that such sentences are not outright contradictory. In suitable contexts, and given suitable backgrounds, being fun may be perceived as bad, while being boring may be perceived as good.<sup>13</sup> Crucially, however, it is not enough to come up with some context in which such a phrase may sound fine in order to conclude that the adjective at stake may be felicitously modified by the phrase "in a good/(bad) way" in the way in which neutral adjectives can. In the case of neutral adjectives, such modification occurs naturally and does not require any special context or stage-setting; not so in the case of evaluative adjectives. This is why it is important to back up the diagnostics of linguistic tests with empirical evidence, and Stojanovic and Kaiser (2022) accomplish this via two routes. First, as already pointed out, they supplement the tests with evidence from psychological norms. Second, they provide naturally occurring instances of sentences such as (11), (12), (19) and (20) drawn from linguistic corpora.

The previous worry concerns linguistic tests more generally. When we are asked to decide whether a given phrase sounds fine or, rather, sounds marked, several problems arise. First, we are often asked to judge the use of a sentence out of any natural context. Second, our intuitions are often unreliable. Third, with enough ingenuity, we can always come up with a context in which the sentence sounds fine. Nevertheless, tests are still a good guide to linguistic typology, provided that they can be fostered with empirical evidence. Empirical data can be obtained experimentally, as well as through the study of linguistic corpora. Experimental studies of evaluative adjectives (broadly construed) include Syrett, Kennedy and Lidz (2010) on gradable adjectives; Liao and Meskin (2018) on extending Syrett's paradigm to aesthetic adjectives; Solt (2016) on gradable adjectives in comparative form; Faroldi and Soria Ruiz (2017) and Soria Ruiz and Faroldi (2021) on extending Solt's paradigm to moral adjectives; Verheyen, Dewin and Egré (2018) on 'heavy' and 'tall'; Smith et al. (2015), Kaiser and Lee (2017), Dinges and Zakkou (2020), Kaiser and Rudin (2020) and Kneer (2021) on predicates of personal taste; Willemsen and Reuter (2021) on moral terms; Ruytenbeek, Verheyen and Spector (2017), Mazzarella and Gotzner (2021) and van Tiel and Pankratz (2021) on the role of adjectival valence (or polarity) in drawing pragmatic inferences. As for (less numerous) corpus studies, let us mention Sassoon (2013) on multidimensional adjectives; McNally and Stojanovic (2017) on aesthetic adjectives; Reuter, Baumgarten and Willemsen (ms.) and McNally and Stojanovic (forthcoming) on moral adjectives.

---

<sup>13</sup> Here is an excerpt, cited by Stojanovic and Kaiser (2022), that provides a felicitous use of "boring in a good way": "Ted was an accountant, and he acted like I expected an accountant to act, which doesn't mean that he was boring, but... Well, let's just say that he was more *adult* than any of my other friends. More *mature*. More... well, yes, boring, but *boring* in a good way" (*The Night We Met*, Rob Byrnes, p. 2; author's italics). Observe, however, how much it takes for the author to build up a context in which the phrase sounds fine.

## 4.2. Two tests for evaluativity based on action guidance

In this and the next section, we shall look at four tests that Soria Ruiz and Stojanovic (2019) propose as a diagnostics for what they call *expressive* (uses of) terms, but what ultimately boils down to *evaluativity* in the sense of section 3.1., understood in terms of value (rather than valence) and contrasted with descriptive and factual language.<sup>14</sup>

Their first two tests are driven by the idea that value-judgments guide action. Recall that one of the tenets of expressivism about value is the observation that by evaluating something positively, one is *ipso facto* motivated to act in its favor, to support it and promote it, while by evaluating something negatively, one is motivated to act against it. This idea can be put to work in order to demarcate evaluative language from descriptive language. In this vein of thinking, Soria Ruiz and Stojanovic (2019) propose a first test that they dub *Juxtaposition with although-type connectives*. In a nutshell, if a statement of the form ‘A thinks that x is *t*, although A has no desire or intention to support or promote x’ is felicitous, and is so in the absence of any specific contextual background, then *t* is a positive evaluative term. Similarly, if a statement of the form ‘A thinks that x is *t*, although A is willing to support or promote x’ is felicitous, and is so in the absence of any specific contextual background, then *t* is a negative evaluative term.

As in the case of the two tests discussed in the previous section, this diagnostics, too, appeals to visible contrasts in perceived felicity of sentences. Compare:

21. Mehmed thinks that providing open access for philosophy journals is a wonderful idea, although he doesn’t have any intention or plan whatsoever of supporting or promoting it.
22. ?? Mehmed thinks that providing open access for philosophy journals is a recent idea, although he doesn’t have any intention or plan whatsoever of supporting or promoting it.

The reason why (22) sounds odd is that a term like "recent" has no evaluative import and does not motivate action in any peculiar way, hence the hearer is at a loss as to the speaker's use of the contrastive *although*. A term like "wonderful", on the other hand, leads to the expectation that one would act in favor of the object to which one would ascribe such a term, which licenses *although* when the expectation is not borne out.

Soria Ruiz and Stojanovic's second test also exploits the idea of action guidance, but in terms of how practical inferences work. They formulate the criterion as follows: "Let  $A_1, \dots, A_n, B \therefore C$  be a cogent practical inference, such that  $A_1, \dots, A_n$  have factual meanings, and none of them is about the agent's practical attitudes or intentions. Then, if B is not about the agent's practical attitudes or intentions either,

---

<sup>14</sup> Soria Ruiz and Stojanovic (2019) opt for the term *expressive*, rather than *evaluative*, because their main aim is to show that expressivism (about evaluative language in general, and moral language in particular) can also be motivated by linguistic evidence. In this respect, they share some of the endeavor of Franzén (2020), even though the latter focuses on the possibility of embeddings under attitude verbs such as 'find', further discussed in section 4.4.

then B has expressive meaning" (Soria Ruiz & Stojanovic 2019, p. 164; again, recall that what they call "expressive meaning" is what we would call "evaluative meaning" along the lines of the Value Approach).

To see how this criterion works in practice, consider a scenario in which an agent is deliberating whether to provide open access to a journal publication. If you start from purely factual premises, such as the premise that providing open access costs a lot of money, or that it makes it possible for a greater number of people to read the journal, such premises by themselves do not settle how to act. However, adding further premises, as in (23) to (26), can bridge this inferential gap:

- 23. Providing open access costs a lot of money. *You don't wish to spend money.*  
∴ Don't provide open access.
- 24. Providing open access makes it likely that many people will read your publication. *You want your publication to be read by many people.*  
∴ Provide open access.
- 25. Providing open access costs a lot of money. *Spending a lot money is a bad idea.* ∴ Don't provide open access.
- 26. Providing open access costs a lot of money. *Providing open access is a generous action.* ∴ Provide open access.

While (23) and (24) involve premises that are explicitly about the agent's attitudes and intentions, (25) and (26) have a surface form of ordinary predicative constructions. However, the practical inferences at stake are cogent. What makes this possible is the presence of evaluative terms, namely 'bad' in (25) and 'generous' in (26).

#### 4.3. Two tests based on epistemic considerations

The second two tests in Soria Ruiz and Stojanovic (2019) rely on the idea that evaluative language does not aim to convey factual information about the world. Rather, it conveys values. This, in turn, is reflected by a difference in how speakers ground their statements. For factual statements, one is warranted to make a statement when one has solid epistemic evidence, such as knowledge. For evaluative statements, no such objective epistemic standards seem to be available. Similarly, conveying false factual information is tantamount to lying, while the very notion of lying seems ill-suited to evaluative language.

These ideas can be cashed out in terms of the following two linguistic tests:<sup>15</sup>

*Lack of epistemic justification:*

---

<sup>15</sup> See Soria Ruiz and Stojanovic (2019), p. 166; we have replaced "expressive" by "evaluative".

Let A be a statement followed by an epistemic evidence questioning device. Then, unless the device is interpreted as questioning the speaker's acquaintance with the object that they are talking about, the exchange is felicitous if A is a factual statement, but markedly less so if A is an evaluative statement.

*Lack of lying potential:*

Factual statements can be lies; evaluative statements cannot.

When put into practice, the two tests rely once again on visible contrasts. Consider a situation in which Mehmed, Zoe and Inés have seen the film *The Wailing* together:

27. Inés: This film is horrible.  
Mehmed: ??How do you know?  
Zoe: ??That's a lie!
28. Inés: This film is censored in China.  
Mehmed: ✓ How do you know?  
Zoe: ✓ That's a lie!

To forestall any misunderstanding, exchanges such "*The Wailing* is a horrible film. - How do you know?" can be felicitous when the question is interpreted as asking how the speaker got to be acquainted with the object that they are evaluating. The speaker might reply, "Because I have seen it" or "Because I have read reviews of it that I trust". Similarly, consider a situation in which Mehmed asks Zoe how she likes his new haircut and she says "It is lovely!" If he replies "You are lying", this may sound felicitous because it will be naturally interpreted as questioning Zoe's sincerity. For these two linguistic tests to do their job, the notions of knowledge and lying must be taken in their literal, epistemic sense.

Soria Ruiz and Stojanovic are also careful to point out that the tests serve to identify evaluative *uses* of terms, rather than evaluative terms (*qua* lexical classes). This brings us back to the idea that evaluativity can apply both at the level of a term's stable lexical meaning and at the level of particular uses in context. This holds equally well for the Valence Approach as the Value Approach. We have pointed out in section 4.1. that terms that have neither a clearly positive nor a clearly negative valence, such as 'surprising' and 'intense', can be, and often are, used to express something positive or negative about the object to which they are ascribed. Thinking now in terms of value scales, rather than valence, we can similarly observe that terms whose primary meanings do not concern values can acquire evaluative meaning in context. Soria Ruiz and Stojanovic illustrate the point with the term "powerful". In its literal sense, the adjective serves to say that the object or individual at stake has power; however, especially in aesthetic discourse, the adjective also serves to convey a positive evaluation. Consider the following exchanges:

29. Inés: Na Hong-jin is powerful.  
Mehmed: (??) How do you know?



Zoe: (??) That's a lie!

Whether or not Mehmed's and Zoe's replies make sense or not will largely depend on how Inés' statement is understood in the first place. If she means to be saying that the director Na Hong-jin has a powerful status in the Korean film industry, then it makes perfect sense to ask her how she knows it, or to accuse her of lying. On the other hand, if she means to be expressing an aesthetic evaluation of Na Hong-jin's cinema, then, just as in the case of (27), the replies sound pretty odd.

#### 4.4. Two tests for subjectivity

In this section, let us look at two criteria that have received a lot of attention in the literature: faultless disagreements and embeddings under subjective attitude verbs like 'find'. Since the notion of evaluativity has not always been carefully developed, it is unclear to which extent these criteria have been taken to apply to evaluative terms, as opposed to subjective terms and to predicates of personal taste. Let us first briefly introduce the ideas, and then see whether they can be used to track evaluativity.

Faultless disagreement is easy to illustrate with disagreements over matters of personal taste. Suppose that Zoe loves oysters while Inés finds their taste and texture disgusting, and compare these two possible dialogues between them:

- 30. Zoe: Oysters are delicious.
- 31. Inés: No, oysters are not delicious.
  
- 32. Zoe: Oysters are high in iron.
- 33. Inés: No, oysters are not high in iron.

Both dialogues appear to be instances of disagreement. However, in (32)-(33), there is an obvious way of settling the matter: given that 100g of raw oysters deliver 4mg of iron, which is considerably more than other kinds of food, Zoe is clearly right while Inés is wrong. In (30)-(31), on the other hand, there is no clear way of settling their disagreement: oysters are delicious to people who like oysters, but not to those who do not like them; *de gustibus non est disputandum*.

There are two issues that we must keep separate. One is the question whether the possibility of a genuine disagreement in which both parties are correct makes sense, and if it does, how to account for it. It may turn out that what looks, at a surface, as a disagreement in which both parties are correct dissolves, upon closer scrutiny, into a case of spurious disagreement. Or it may turn out that it is a genuine disagreement, but one for which one of the disagreeing parties turns out to be wrong, even if it is unclear how to adjudicate who is wrong. What is going on in dialogues such as (30)-(31) can be understood in many different ways; for discussion of various strategies to account for the phenomenon, see e.g. Stojanovic (2017), Karczewska (2019), Zakkou (2019), Zeman (2020), and the numerous references therein.

The second issue, which is the one that interests us here, is how this diagnostics carves up the linguistic space. In other words, we wish to know which terms pattern like 'delicious' in (30)-(31), and which pattern like 'high in iron' in (32)-(33). If natural language expressions split neatly into two such classes, then faultless disagreement can indeed be used as a linguistic criterion to classify them accordingly.

The problem, however, is that we are facing a very versatile phenomenon. Even the expression that we have previously contrasted with 'delicious', namely, 'high in iron', can easily give rise to disagreements that cannot be clearly settled. Compare (32)-(33) with the following:

- 34. Zoe: Green beans are high in iron.
- 35. Mehmed: No, they are not.

Zoe and Mehmed can be equally knowledgeable about the fact that 100g of green beans deliver 1mg of iron. Still, their disagreement may be driven by a difference in what they take to be the relevant standards for some food to count as "high" in iron. Mehmed's standards may be more demanding than Zoe's, so that only those foods that deliver, say, over 3mg per 100g count as sufficiently high in iron.<sup>16</sup> In general, any term that allows for some vagueness, or that lends itself to different interpretations, can, in suitable contexts, generate disagreements that cannot be settled by appeal to facts; that is, *faultless* disagreements (see e.g. Sundell 2011, Odrowąż-Sypniewska 2021).

Here it may help to recall a point that we stressed in section 3.2. Namely, when it comes to adjectives, it is often insightful to see how they behave in the comparative:

- 36. Zoe: Oysters are tastier than green beans.
- 37. Inés: No, green beans are tastier.
  
- 38. Zoe: Oysters are higher in iron than green beans.
- 39. Inés: No, green beans as high in iron as oysters.

While the dialogue in (36)-(37) appears as faultless as the one in (30)-(31), the one is (38)-(39) can be definitely settled by facts. Given that oysters deliver four times more iron than green beans, Zoe is right, Inés is wrong. The upshot, then, is that even if faultless disagreement turns out to be a very versatile phenomenon, it is still a helpful tool in identifying an interesting class of adjectives. But is the class at stake that of evaluative adjectives, or some other class?

The way in which faultless disagreement has been used in the literature is as a diagnostics to identify subjectivity (see Willer (this volume)), and, more specifically,

---

<sup>16</sup> One could wonder why (32)-(33) is, then, a case in which the disagreement *can* be clearly settled. The reason is, roughly, that for Inés to justify reference to standards that are so extremely stringent that even oysters fail to meet them ends up completely unreasonable. This squares nicely with the proposal defended in Vardomskaya (2018), who analyzes the objective-subjective divide in terms of whether there is a reliable, socially enforced consensus as to what evidence matters to determine truth and/or settle disagreements.

predicates of personal tastes, or PPTs, for short (see Bylinina 2014, Umbach 2021). Do evaluative adjectives behave in the same way? Note that the PPTs that are most often discussed are 'delicious', 'tasty', 'disgusting', 'fun' and 'boring'. All five are evaluative, both in the Valence sense and in the Value sense. 'Delicious', 'tasty' and 'fun' have positive valence, 'disgusting' and 'boring' negative. 'Delicious' and 'tasty' are used to place items on the higher end of a *gustatory value* scale, 'disgusting' (when used for food) on its lower end.; 'fun' and 'boring' have analogous functions with respect to what one might dub a *ludicity value* scale.<sup>17</sup>

The question, then, is whether all evaluative terms necessarily give rise to faultless disagreement. We submit that this is far from obvious. To be sure, the answer ultimately depends both on how one understands faultless disagreement and how one approaches the notion of evaluativity. But to give at least a sense for why we think that the answer is likely to be negative, take the case of moral terms, as in:

40. Zoe: Female Genital Mutilation is unethical.

41. Inés: No, it isn't.

Female Genital Mutilation (FGM) is an extremely harmful practice internationally recognized as a violation of the human rights of girls and women, and it has been banned by the UN since 2012. As the World Health Organization puts it, "the practice also violates a person's rights to health, security and physical integrity, the right to be free from torture and cruel, inhuman or degrading treatment, and the right to life when the procedure results in death" (<https://www.who.int/news-room/fact-sheets/detail/female-genital-mutilation>). Now, suppose that Inés thinks that since there are societies that still endorse this practice, this must mean that FGM is not unethical after all. Nevertheless, hardly anyone will be moved by such arguments, nor would we want to say that the disagreement in (40)-(41) is faultless. While we are all fairly happy to accept that oysters are delicious to those who like them, but not to those who do not like them, we would balk away from the idea that FGM is ethical to those societies that endorse it, and unethical to those that do not. When it comes to moral values, disagreements tend to be more robust, and to resist relativization to a subjective point of view.

To be sure, the observation that moral disagreements do not trigger the intuition of faultlessness to the same degree as disagreements over matters of taste do does not mean that in the moral case, the disagreement can be settled merely by looking at the facts. What grounds judgments of morality, and how disagreements over moral values are to be understood, are longstanding questions in metaethics and in moral theory. The point is, rather, that if faultless disagreement is to be used as a linguistic test, then moral terms will not pattern quite like subjective terms such as PPTs. True enough, they may not pattern like descriptive, factual terms either. Indeed, psychological

---

<sup>17</sup> It is actually more accurate to assume that, in each case, there are two distinct scales rather than a single scale; namely, a "tastes good"-scale, which ranks food as to how good it tastes (tasty, delicious, worthy of one, two, three Michelin stars) and a "tastes bad"-scale; similarly for other domains (see e.g. Israel 2004, p. 703). Since we shall not discuss scalar inferences and implicatures, nothing important hinges on whether we talk of a single or separate scales.

studies that investigate folk intuitions regarding moral disagreements (see Goodwin and Darley (2008) for one of the first such studies, and Sarkissian (2016) for an overview), tell us that morality lies midway between clearly subjective issues, such as judgments of personal taste, and clearly objective issues, such as factual matters. But, to bring the point home, faultless disagreement primarily tracks subjectivity (see Willer (this volume); see also Kompa (2016) and Rudolph (2019), whose proposals arguably go in a similar direction). Even if certain evaluative terms are prone to elicit faultless disagreement, they do so in virtue of being subjective terms, too: hence we should not mistake the phenomenon for a linguistic criterion for evaluativity.

Let us close this section by gesturing toward another popular test: embeddability under subjective attitude verbs such as 'find'. Again, though, we submit that this test, too, tracks subjectivity rather than evaluativity. Consider the following contrasts:

- 42. Zoe: I find this film superb / boring / intriguing / powerful.
- 43. Inés: # I find this film Korean / censored in China / 16-rated.

The adjectives used by Zoe in (42) can all be felicitously embedded under 'find'. Zoe's statement reflects her view that the film is superb, boring, interesting or powerful. By contrast, adjectives such as 'Korean', 'censored' or '16-rated' cannot be embedded under such attitude verbs; (43) cannot be used to report Inés' belief that the film is Korean, censored in China, or that it is rated as unsuitable for viewers aged below 16.

Embeddability under 'find' was first used in Sæbø (2009) as a linguistic criterion to delineate PPTs. But he also used it to argue that PPTs are sensitive to an implicit experiencer argument; that is to say, that they encode in their semantics an implicit reference to one or more subjects who have direct experience of the object to which the expression is ascribed. The subsequent literature disagrees whether adjectives that are felicitous under 'find' need to be experiencer-sensitive or not. The first camp included, in addition to Sæbø, Bylinina (2014, 2017), McNally and Stojanovic (2017), and Stojanovic (2019b). The second includes Kennedy (2013), Kennedy and Willer (2016), Umbach (2017), Coppock (2018), Vardomskaya (2018), Franzén (2020), Silk (2021), and Anand and Korotkova (2021).

The question for us is, once again, whether embeddability under 'find' can be seen as a test that identifies evaluative adjectives. Note that the adjectives in (42) appear to be evaluative; more precisely, 'superb' and 'boring' are evaluative both in the Valence sense and in the Value sense, while 'intriguing' and 'powerful' do not come with a lexically encoded valence, but can be used to express evaluative judgments in context. However, as before, we suggest that when evaluative adjectives are felicitous under 'find', this is because the adjectives under consideration are also subjective. In other words, the test tracks subjectivity rather than evaluativity, but since there is a substantive overlap between the two classes (on both understandings of evaluativity), many evaluative adjectives turn out to be embeddable under subjective attitude verbs such as 'find'.

Why think, then, that 'find' does not track evaluativity? Take again the case of moral adjectives, and compare:

44. Zoe finds Na Hong-jin's films boring.  
 45. ? Zoe finds Female Genital Mutilation unethical.

While (45) need not sound outright infelicitous, and, at any rate, sounds less so than (43), it is an odd way of reporting Zoe's moral judgment that FGM is unethical, and sounds much less natural than (44). The reason, we submit, is that 'finds' introduces an experiential component, and while it is true that some evaluative judgments are based on one's first-hand experience of the object, other evaluative judgments are not. In (44), Zoe's judgment that the film is boring implies that she must have watched it (at least some of it). To make sense of (45), we tend to reconstruct the context as one in which Zoe's judgment that FGM is unethical is somehow based on her experience.

To wrap up, the 'find' test, despite its popularity, yields data that tend to be murky when we move away from clearly subjective adjectives. PPTs, and adjectives that involve experiencers, are subjective and therefore clearly felicitous under 'find'. But among evaluative adjectives that are not clearly subjective, and moral adjectives in particular, 'find' is not a preferred choice to report a value judgment.<sup>18</sup>

## 5 Conclusion

Contemporary philosophy of language has witnessed an outburst of interest in the topic of evaluative meaning. At the same time, the issue of what it takes for an expression, statement or judgment to be evaluative has been relatively underexplored. The problem is that there is no single, widely agreed upon pretheoretic notion of evaluativity. True enough, there is agreement on a core class of cases that are to be included in the realm of the evaluative: judgments about what is good or bad, what is beautiful or ugly, what is just or unjust. But even for the core cases, when we inquire deeper into their evaluative nature, two views of evaluativity appear to compete with each other. The first view, which we have called the Valence Approach, takes such judgments to be evaluative because they convey something positive or negative about the object of evaluation. Relatedly, it identifies words such as 'beautiful', 'ugly', 'good' and 'bad' as evaluative predicates because they are systematically used to express such positive or negative judgments. We have also pointed out that expressivist approaches to value judgments fit well with the Valence Approach because they operate with an analogous distinction at the level of attitudes: *pro* attitudes correlate with positive value judgments, and *con* attitudes, with negative ones. The second view, that is, the Value Approach, ties the notion of evaluativity to the idea of assignment of value. While nonevaluative statements express factual information, evaluative statements

---

<sup>18</sup> It bears mentioning that the linguistic data have not been sufficiently studied from an empirical point of view, and that, additionally, there are important cross-linguistic variations between 'find' in English and its near-equivalents in other languages, like 'trouver' in French, 'finden' in German and 'tycka' in Swedish. The claim that 'find' is not a preferred choice when it comes to moral adjectives finds empirical support in the recent corpus study in Stojanovic and McNally (forthcoming).

reflect how the object of evaluation ranks on some corresponding value scale.

Our modest aim in this chapter has been to articulate the distinction between the two approaches, and show how depending on which approach one endorses, different, albeit partially overlapping classes of expressions get to be identified as "evaluative". Now, to complicate the matter even further, the topic of evaluative meaning is most often discussed in parallel with the topic of *subjective* meaning. How evaluativity and subjectivity relate to each other, though, remains an open philosophical issue.

Let us close by pointing out two limitations of what we have covered in this chapter. First, evaluativity has been traditionally linked to the topic of *normativity*. Judgments about what is good or bad are often taken not only to express values, but also to issue norms and prescriptions. Statements such as "It is bad to lie" or "Lying is wrong" are often seen as normative and take the force of an imperative: *Don't lie!* Conversely, to say that it is good to do something may be a way of saying that one ought to do it. While evaluative expressions, especially 'good' and 'bad', seem to be related to deontic expressions such as 'ought', the exact nature of their relationship is still under study; see e.g. Tappolet (2013), Lassiter (2017), and Soria Ruiz, Maldonado and Stojanovic (2022) for discussion.

Second, the class of evaluative terms on which we have focused are adjectives. The choice is understandable, given that adjectives are the most frequent carriers of evaluativity, both in the Valence and in the Value sense. But evaluative meaning is definitely not restricted to adjectives. In particular, there is a thriving literature on certain *nouns* that carry negative evaluations. The case at point are pejoratives, which include racial and other slurs, such as the N-word, and insults, such as 'jerk' or 'asshole'.<sup>19</sup> We have deliberately abstained from discussing them, since that would have taken us far beyond the scope of the present chapter.<sup>20</sup>

## Acknowledgments

I would like to thank Ernie Lepore and Una Stojnic for the invitation to contribute to this handbook. My thoughts on this topic have been shaped over the years thanks to many fruitful discussions with colleagues from the Jean-Nicod Institute in Paris, the GLiF and LOGOS groups in Barcelona, and elsewhere. I am particularly grateful to Bianca Cepollaro, Andrés Soria Ruiz, Louise McNally, Teresa Marques, Elsi Kaiser, Julia Zakkou, Filippo Domaneschi, Neftalí Villanueva, Dan Zeman, Markus Kneer and Nils Franzén. I would also like to acknowledge institutional support from the project ANR-17-EURE-0017 FrontCog.

---

<sup>19</sup> For the differences between the two types of derogatory terms, both in terms of how they are identified and in terms of how they behave vis-à-vis other terms, see e.g. Diaz-Legaspe (2019), Cepollaro, Sulpizio and Bianchi (2019), Cepollaro, Domaneschi and Stojanovic (2020), Jeshion (2021), Stojanovic (2021).

<sup>20</sup> One may wonder how much there is in common between evaluativity in the adjectival realm and the sort of evaluative meaning that derogatory terms instantiate. Cepollaro and Stojanovic (2016) and Cepollaro (2020) argue that slurs and thick evaluative adjectives, like 'generous' and 'cruel', call for a unified semantic treatment.

## References

- Anand, Pranav and Natasha Korotkova (2021) How to theorize about subjective meaning: a lesson from 'de re'. *Linguistics and Philosophy*. <https://doi.org/10.1007/s10988-021-09331-0>
- Bedke, Matthew (2017) Cognitivism and non-cognitivism. In Tristram McPherson & David Plunkett (eds.), *The Routledge Handbook of metaethics*. Routledge. pp. 292-307.
- Bianchi, Claudia (2014) Slurs and appropriation: an echoic account. *Journal of Pragmatics* 66: 35--44.
- Bierwisch, Manfred (1989) The Semantics of gradation, in Bierwisch, M. and Ewald Lang (eds.), *Dimensional Adjectives*. Berlin: Springer-Verlag, pp. 71-261.
- Braşoveanu, Adrian and Rett, Jessica (2017) Evaluativity across adjective and construction types: An experimental study. *Journal of Linguistics* 54: 263-329.
- Bylinina, Lisa (2014) *The grammar of standards: judge-dependence, purpose-relativity, and comparison classes in degree constructions*. LOT Dissertation Series 347. Utrecht: LOT.
- Bylinina, Lisa (2017) Judge-dependence in degree constructions. *Journal of Semantics* 34: 291-331.
- Camp, Elisabeth (2017) Metaethical expressivism. In Tristram McPherson & David Plunkett (eds.), *The Routledge Handbook of metaethics*. Routledge. pp. 87-101.
- Cepollaro, Bianca (2017) When evaluation changes. An echoic account of appropriation and variability. *Journal of Pragmatics* 117: 29–40.
- Cepollaro, Bianca (2020) *Slurs and thick terms*. Rowman & Littlefield.
- Cepollaro, Bianca and Isidora Stojanovic (2016) Hybrid evaluatives: in defense of a presuppositional account. *Grazer Philosophische Studien* 93: 458-488.
- Cepollaro, Bianca, Simone Sulpizio and Claudia Bianchi (2019) How bad is it to report a slur? An empirical investigation. *Journal of Pragmatics* 146: 32-42
- Cepollaro, Bianca, Filippo Domaneschi and Isidora Stojanovic (2020) When is it ok to call someone a jerk? An experimental investigation of expressives. *Synthese*. <https://doi.org/10.1007/s11229-020-02633-z>
- Contesi, Filippo (2018) Art and emotion. *Oxford Bibliographies*. <https://doi.org/10.1093/OBO/9780195396577-0381>.
- Coppock, Elizabeth (2018) Outlook-based semantics. *Linguistics and Philosophy* 41: 125–164.
- Diaz-Legaspe, Justina (2019) What is a slur? *Philosophical Studies*. DOI: 10.1007/

s11098- 019-01259-3.

- Dinges, Alexander and Julia Zakkou (2020) A direction effect on taste predicates. *Philosophers' Imprint* 20:1-22.
- Faroldi, Federico and Andrés Soria Ruiz (2017) The Scale structure of moral adjectives. *Studia Semiotyczne* 31: 161-178.
- Franzén, Nils (2018) Aesthetic evaluation and first-hand experience. *Australasian Journal of Philosophy* 96.
- Franzén, Nils (2020) Evaluative discourse and affective states of mind. *Mind* 129: 1095–1126.
- Goodwin, Geoffrey and John Darley (2008) The psychology of meta-ethics: exploring objectivism. *Cognition* 106: 1339–1366.
- Israel, Michael (2004) The Pragmatics of polarity. In Horn, Laurence and Ward, Gregory (eds.) *The Handbook of Pragmatics*, Blackwell. pp. 701-723.
- Jeshion, Robin (2021) Varieties of pejoratives. In Khoo, Justin and Rachel Sterken (eds.) *The Routledge Handbook of social and political philosophy of language*.
- Kaiser, Elsi and Lee, Jamie Herron (2017) Experience matters: a psycholinguistic investigations of predicates of personal taste. *Proceedings of SALT-27*.
- Kaiser, Elsi & Rudin, Deniz (2020). When faultless disagreement is not so faultless: What widely-held opinions can tell us about subjective adjectives. In P. Farrell (ed.), *Proceedings of the Linguistic Society of America* 5, 698-707.
- Karczewska, Natalia (2019) *Faultless disagreement in contemporary semantic theories*. PhD Dissertation. University of Warsaw.
- Kennedy, Christopher (2013) Two sources of subjectivity: Qualitative assessment and dimensional uncertainty. *Inquiry* 56. 258–277.
- Kennedy, Christopher and Malte Willer (2016) Subjective attitudes and counterstance contingency. *Proceedings of SALT* 26.
- Kompa, Nikola (2016) Faultless disagreement, context-sensitivity and the semantics of evaluation. *Grazer Philosophische Studien* 93: 396-418.
- Kneer, Markus (2021) Predicates of personal taste: empirical data. *Synthese*. <https://doi.org/10.1007/s11229-021-03077-9>
- Lassiter, Dan (2017). *Graded modality: qualitative and quantitative perspectives*. Oxford University Press.
- Liao, Shen-yi and Aaron Meskin (2017) Aesthetic adjectives: experimental semantics and context-sensitivity. *Philosophy and Phenomenological Research* 94: 371–398.
- Liao, Shen-yi, Louise McNally and Aaron Meskin (2016) Aesthetic adjectives lack



- uniform behavior. *Inquiry* 59: 618-631.
- Mazzarella, Diana and Nicole Gotzner (2021) The polarity asymmetry of negative strengthening: dissociating adjectival polarity from face-threatening potential. *Glossa: a journal of general linguistics* 6(1): 47. 1–17. DOI: <https://doi.org/10.5334/gjgl.1342>
- McNally, Louise and Isidora Stojanovic (2017) Aesthetic adjectives, in Young, James (ed.) *The semantics of aesthetic judgment*. Oxford: OUP. pp. 17-37.
- Mohammad, Saif M. (2018) Obtaining reliable human ratings of valence, arousal, and dominance for 20,000 English words. *Proceedings of the Annual Conference of the Association for Computational Linguistics (ACL)*. Melbourne, Australia.
- Odrowąż-Sypniewska, Joanna (2021) Vagueness in natural language. In *The Cambridge handbook of the philosophy of language*, ed. Piotr Stalmaszczyk, pp. 434–449. Cambridge: Cambridge University Press.
- Reuter, Kevin, Lucien Baumgarten and Pascale Willemsen (ms.) Tracing Thick and Thin Concepts Through Corpora.
- Rudolph, Rachel (2021) Talking about appearances: the roles of evaluation and experience in disagreement. *Philosophical Studies* 177: 197-217.
- Ruytenbeek, Nicolas, Steven Verheyen & Benjamin Spector (2017) Asymmetric inference towards the antonym: Experiments into the polarity and morphology of negated adjectives. *Glossa: A Journal of General Linguistics* 2(1). 92. DOI: <https://doi.org/10.5334/gjgl.151>
- Sarkissian, Hagop. 2016. Aspects of folk morality: Objectivism and relativism. In *A companion to experimental philosophy*, ed. Wesley Buckwalter and Justin Systema, 212– 224. Malden and Oxford: Wiley Blackwell. <https://doi.org/10.1002/9781118661666.ch14>
- Sæbø, Kjell Johan (2009) Judgement ascriptions. *Linguistics and Philosophy* 32: 327-352.
- Silk, Alex (2021) Evaluational adjectives. *Philosophy and Phenomenological Research* 102: 127-161.
- Sassoon, Galit W. (2013) A typology of multidimensional adjectives. *Journal of Semantics* 30: 335-380.
- Solt, Stephanie (2016) Ordering subjectivity and the relative-absolute distinction. *Proceedings of Sinn und Bedeutung* 20.
- Soria Ruiz, Andrés (2019) *The Place of value in natural language*. PhD Dissertation, Ecole Normale Supérieure, Paris.
- Soria Ruiz, Andrés and Isidora Stojanovic (2019) On Linguistic evidence for expressivism. *Philosophy* 86: 155-180.

- Soria Ruiz, Andrés and Federico Faroldi (2020) Moral adjectives, judge-dependency and holistic multidimensionality. *Inquiry* 64: 1-30.
- Soria Ruiz, Andrés, Bianca Cepollaro and Isidora Stojanovic (2021) The Semantics and pragmatics of value judgments. In Stalmaszczyk, Piotr (Ed.) *The Cambridge Handbook of the Philosophy of Language*, Cambridge University Press, pp. 434-449.
- Soria Ruiz, Andrés, Mora Maldonato and Isidora Stojanovic (2022) 'Good' and 'ought' in argumentation: COVID-19 as a case study. In Oswald, S., M. Lewiński, S. Greco & S. Villata (Eds.) *The Pandemic of Argumentation*. Springer. pp. 43-64.
- Smith, Elizabeth Allyn, Elena Castroviejo, Laia Mayol and Monique Beaudry (2015) Cross-linguistic experimental evidence distinguishing the role of context in disputes over taste and possibility. In Christiansen, Henning and Isidora Stojanovic (Eds.) *Modeling and Using Context. Proceedings of The Ninth International and Interdisciplinary Conference CONTEXT 2015*. Springer Lecture Notes in Artificial Intelligence.
- Stojanovic, Isidora (2016) Expressing aesthetic judgments in context. *Inquiry* 59: 663-685.
- Stojanovic, Isidora (2017) Context and disagreement. *Cadernos de Estudos Lingüísticos* 59: 9-22.
- Stojanovic, Isidora (2019a) Disagreements about taste vs. disagreements about moral issues. *American Philosophical Quarterly* 56 (2019): 29-41.
- Stojanovic, Isidora (2019b) An Empirical approach to aesthetic adjectives. In Cova, Florian and Sébastien Réhault (Eds.), *Advances in experimental philosophy of aesthetics*. Oxford: Bloomsbury Academic. 2018. 221-240.
- Stojanovic, Isidora (2021) Derogatory terms in free indirect discourse, in Maier, Emar & Stokke, Andreas (Eds.) *The Language of Fiction*. Oxford University Press. pp. 349-378.
- Stojanovic, Isidora and Elsi Kaiser (2022) Exploring valence in judgments of taste. In Wyatt, Jeremy, Julia Zakkou and Dan Zeman (Eds.) *Perspectives on Taste*. Routledge. pp. 231-259.
- Stojanovic, Isidora and Louise McNally (forthcoming) Are moral predicates subjective? A corpus study, in Bordonaba, David (ed.) *Experimental Philosophy of Language: Perspectives, Methods and Prospects*. Springer.
- Sundell, Timothy (2011) Disagreements about taste. *Philosophical Studies* 155: 267-288.
- Syrett, Kristen, Chris Kennedy and Jeffrey Lidz (2010) Meaning and context in children's understanding of gradable adjectives. *Journal of Semantics* 27: 1-35.
- Tappolet, Christine (2013) Evaluative vs. deontic concepts. In Lafollette, Hugh

- (ed.), *International Encyclopedia of Ethics*. Wiley-Blackwell. pp. 1791-99.
- Umbach, Carla (2017) Evaluative propositions and subjective judgments, in J. van Wijnbergen-Huitink & C. Meier (eds.) *Subjective meaning*. Berlin: De Gruyter.
- Umbach, Carla (2021) Evaluative predicates beyond 'fun' and 'tasty'. In Daniel Gutzmann, Lisa Matthewson, Cecile Meier, Hotze Rullmann, Thomas Ede Zimmermann (eds.) *The Wiley Blackwell Companion to Semantics*.
- van Roojen, Mark (2018) Moral cognitivism vs. non-cognitivism. In Zalta, Edward (ed.) *The Stanford Encyclopedia of Philosophy* (Fall 2018 Edition).
- van Tiel, Bob and Elizabeth Pankratz (2021) Adjectival polarity and the processing of scalar inferences. *Glossa: a journal of general linguistics* 6(1): 32. 1–21. DOI: <https://doi.org/10.5334/gjgl.1457>
- Väyrynen, Pekka (2013) *The rude, the lewd and the nasty*. Oxford University Press.
- Verheyen, Steven, Sabrina Dewil and Paul Egré (2018) Subjectivity in gradable adjectives: The case of *tall* and *heavy*. *Mind and Language*. <https://doi.org/10.1111/mila.12184>
- Warriner, A. B., Kuperman, V., & Brysbaert, M. (2013) Norms of valence, arousal, and dominance for 13,915 English lemmas. *Behavior Research Methods*, 45(4).
- Willemsen, Pascale and Reuter, Kevin (2021) Separating the evaluative from the descriptive: an empirical study of thick concepts. *Thought*.
- Willer, Malte (this volume) Subjectivity. In Lepore, Ernie and Una Stojnic (Eds.) *The Oxford Handbook of Contemporary Philosophy of Language*.
- Zakkou, Julia (2019) Faultless disagreement: A defense of contextualism in the realm of personal taste. Klostermann.
- Zeman, Dan (2020) Faultless disagreement. In Martin Kusch (ed.), *The Routledge Handbook of Philosophy of Relativism*. Routledge. pp. 486-495.